

hw1

October 18, 2023

```
[ ]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import sklearn.model_selection
import sklearn.preprocessing
```

0.1 Preprocessing Work

```
[ ]: data = pd.read_excel("Concrete_Data.xls")
shortenedCols = ["Cement", "Slag", "Ash", "Water", "Superplasticizer", "Age", "Strength"]
newCols = dict(zip(data.columns, shortenedCols))

renamed = data.rename(columns=newCols)
data.head()
```

```
[ ]: Cement (component 1)(kg in a m^3 mixture) \
0          540.0
1          540.0
2          332.5
3          332.5
4          198.6

Blast Furnace Slag (component 2)(kg in a m^3 mixture) \
0          0.0
1          0.0
2          142.5
3          142.5
4          132.4

Fly Ash (component 3)(kg in a m^3 mixture) \
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0
```

	Water (component 4)(kg in a m ³ mixture) \	
0	162.0	
1	162.0	
2	228.0	
3	228.0	
4	192.0	

	Superplasticizer (component 5)(kg in a m ³ mixture) \	
0	2.5	
1	2.5	
2	0.0	
3	0.0	
4	0.0	

	Coarse Aggregate (component 6)(kg in a m ³ mixture) \	
0	1040.0	
1	1055.0	
2	932.0	
3	932.0	
4	978.4	

	Fine Aggregate (component 7)(kg in a m ³ mixture)	Age (day) \
0	676.0	28
1	676.0	28
2	594.0	270
3	594.0	365
4	825.5	360

	Concrete compressive strength(MPa, megapascals)
0	79.986111
1	61.887366
2	40.269535
3	41.052780
4	44.296075

```
[ ]: train_size = 900/renamed.shape[0]
assert(train_size == 900/1030)

y = renamed["Strength"]
X = renamed.drop(columns="Strength")

X_train, X_test, y_train, y_test = sklearn.model_selection.train_test_split(X,
    ↪y, train_size=train_size)
X_train = np.array(X_train)
X_test = np.array(X_test)
y_train = np.array(y_train)
y_test = np.array(y_test)
```

```
[ ]: scaler = sklearn.preprocessing.StandardScaler()
X_scaler = scaler.fit(X_train)

X_train_norm = np.array(X_scaler.transform(X_train))
X_test_norm = np.array(X_scaler.transform(X_test))

y_scaler = sklearn.preprocessing.StandardScaler()
y_scaler.fit(y_train.reshape(-1, 1))
y_train_norm = y_scaler.transform(y_train.reshape(-1, 1)).ravel()
y_test_norm = y_scaler.transform(y_test.reshape(-1, 1)).ravel()
```

```
[ ]: # This function was used to generate the plots in the report
# I have disabled its execution to save space in the final report

def feature_dist(df):
    for feature in shortenedCols:
        plt.figure(figsize=(3, 2))
        plt.hist(df[feature], bins=50)
        plt.title(f'Distribution of {feature}')
        plt.xlabel(feature)
        plt.ylabel('Frequency')
        plt.show()

# feature_dist(renamed)
```

0.2 Gradient Descent Algorithm

```
[ ]: def gradient_descent(X, y, learning_rate, iterations):
    X = np.c_[np.ones(X.shape[0]), X]
    theta = np.zeros(X.shape[1])
    for _ in range(iterations):
        error = np.dot(X, theta.T) - y
        gradient = (2 / len(y)) * np.dot(X.T, error)
        theta = theta - (learning_rate * gradient)

    return theta
```

0.2.1 Univariate Model

For each feature, we calculate the unnormalized and normalized theta values. Before running the gradient descent algorithm, we concatenate a column of 1s before our feature to account for the “constant” offset term. Therefore, we get two theta values back per run. The first corresponds to the constant, and the second corresponds to the scale of the first feature.

It’s important to note the difference in learning rate between the unnormalized and normalized values. Because of the larger unnormalized values, running the algorithm with a larger learning rate will cause our numbers to overflow. This is partially accounted for by increasing the number

of iterations.

```
[ ]: thetas = np.zeros([8,2])
norm_thetas = np.zeros([8, 2])
for i in range(8):
    target = np.array(X_train[:, i]).reshape(-1,1)
    thetas[i] = gradient_descent(target, y_train, 0.000001, 1000)

    norm_target = np.array(X_train_norm[:, i]).reshape(-1,1)
    norm_thetas[i] = gradient_descent(norm_target, y_train_norm, 0.005, 10000)

print(thetas)
print(norm_thetas)
```

```
[[0.0034637  0.12030553]
 [0.03997999 0.21899004]
 [0.04575242 0.25362762]
 [0.00310929 0.19011898]
 [0.06798608 0.47597249]
 [0.00096533 0.03601923]
 [0.00141427 0.04509721]
 [0.04412275 0.3226835 ]]
[[ 5.55308885e-17  5.02343218e-01]
 [ 1.38733469e-16  1.46692528e-01]
 [ 1.28356585e-16 -1.12944599e-01]
 [-3.46991571e-16 -2.87010654e-01]
 [ 1.52643330e-16  3.66130334e-01]
 [ 2.86180955e-16 -1.71369193e-01]
 [ 2.49810049e-16 -1.75216548e-01]
 [ 1.21386851e-16  3.30158472e-01]]
```

0.3 Multivariate Model

This uses the exact same function defined above, yet the entire input is passed to the gradient descent algorithm instead of just a single feature. We expect the output of each function to be a 1x9 array, accounting for the extra constant column of 1s.

```
[ ]: multivar_thetas = gradient_descent(X_train, y_train, 0.0000001, 10000)
multivar_norm_thetas = gradient_descent(X_train_norm, y_train_norm, 0.01, 1000)

print(multivar_thetas)
print(multivar_norm_thetas)
```

```
[-4.94565987e-05  1.17061494e-01  9.58321104e-02  8.95407190e-02
 -1.34837471e-01  3.34834230e-02  7.90934998e-04  1.21552682e-02
 1.04428152e-01]
[-3.87443164e-16  5.73305067e-01  3.59729574e-01  1.79375515e-01
 -3.32187173e-01  8.88597157e-02 -5.12532722e-02 -7.57791740e-02
 4.21943592e-01]
```

0.4 Training Model Performance

```
[ ]: def predict(X, thetas):
    preds = np.dot(X, thetas.T)
    return preds

def mse(y, y_pred):
    return np.sum(np.square(y - y_pred)) / np.size(y)

[ ]: print("Performance on Training Dataset")
print("RSquared\t\t MSE\t\t Variance")

X_multivar_train = np.c_[np.ones(X_train.shape[0]), X_train]
y_pred = predict(X_multivar_train, multivar_thetas)
MSE = mse(y_train, y_pred)
var = np.var(y_train)
rsquared = 1 - (MSE/var)
print(f"{rsquared}, {MSE}, {var} = Unnormalized Multivariate R^2")

X_multivar_norm_train = np.c_[np.ones(X_train_norm.shape[0]), X_train_norm]
y_pred = predict(X_multivar_norm_train, multivar_norm_thetas)
MSE = mse(y_train_norm, y_pred)
var = np.var(y_train_norm)
rsquared = 1 - (MSE/var)
print(f"{rsquared}, {MSE}, {var} = Normalized Multivariate R^2\n")

for i in range(8):
    target = np.c_[np.ones((X_train.shape[0], 1)), X_train[:,i]]
    y_pred = predict(target, thetas[i])
    MSE = mse(y_train, y_pred)
    var = np.var(y_train)
    rsquared = 1 - (MSE/var)
    print(f"{rsquared}, {MSE}, {var} = R^2 for Unnormalized Feature_{shortenedCols[i]}")

print("")

for i in range(8):
    target = np.c_[np.ones((X_train_norm.shape[0], 1)), X_train_norm[:,i]]
    y_pred = predict(target, norm_thetas[i])
    MSE = mse(y_train_norm, y_pred)
    var = np.var(y_train_norm)
    rsquared = 1 - (MSE/var)
    print(f"{rsquared}, {MSE}, {var} = R^2 for Normalized Feature_{shortenedCols[i]}")
```

Performance on Training Dataset

RSquared	MSE	Variance	
0.6056716878185033,	110.91766254069823,	281.28252299988645	= Unnormalized Multivariate R ²
0.6143422909001132,	0.38565770909988684,	1.0	= Normalized Multivariate R ²
0.18211356369270126,	230.05716033190294,	281.28252299988645	= R ² for Unnormalized Feature Cement
-2.2763111974318657,	921.5690797464143,	281.28252299988645	= R ² for Unnormalized Feature Slag
-2.87211977997769,	1089.15962106989,	281.28252299988645	= R ² for Unnormalized Feature Ash
-0.20024187017790296,	337.60706145374274,	281.28252299988645	= R ² for Unnormalized Feature Water
-3.664915222607407,	1312.1591233955883,	281.28252299988645	= R ² for Unnormalized Feature Superplasticizer
-0.08772754903166602,	305.9587493281097,	281.28252299988645	= R ² for Unnormalized Feature Coarse
-0.12338612287634709,	315.98888294571935,	281.28252299988645	= R ² for Unnormalized Feature Fine
-2.2268732003673697,	907.6630352000519,	281.28252299988645	= R ² for Unnormalized Feature Age
0.25234870854805735,	0.7476512914519426,	1.0	= R ² for Normalized Feature Cement
0.02151869785031757,	0.9784813021496824,	1.0	= R ² for Normalized Feature Slag
0.012756482449876483,	0.9872435175501235,	1.0	= R ² for Normalized Feature Ash
0.08237511534942799,	0.917624884650572,	1.0	= R ² for Normalized Feature Water
0.13405142172356965,	0.8659485782764303,	1.0	= R ² for Normalized Feature Superplasticizer
0.029367400160310808,	0.9706325998396892,	1.0	= R ² for Normalized Feature Coarse
0.030700838713056733,	0.9692991612869433,	1.0	= R ² for Normalized Feature Fine
0.10900461661907535,	0.8909953833809247,	1.0	= R ² for Normalized Feature Age

0.5 Testing Model Performance

```
[ ]: print("Performance on Testing Dataset")
      print("RSquared\t\t MSE\t\t Variance")

X_multivar_test = np.c_[np.ones(X_test.shape[0]), X_test]
y_pred = predict(X_multivar_test, multivar_thetas)
MSE = mse(y_test, y_pred)
var = np.var(y_test)
rsquared = 1 - (MSE/var)
print(f"{rsquared}, {MSE}, {var} = Unnormalized Multivariate R^2")

X_multivar_norm_test = np.c_[np.ones((X_test_norm.shape[0], 1)), X_test_norm]
y_pred = predict(X_multivar_norm_test, multivar_norm_thetas)
```

```

MSE = mse(y_test_norm, y_pred)
var = np.var(y_test_norm)
rsquared = 1 - (MSE/var)
print(f"{rsquared}, {MSE}, {var} = Normalized Multivariate R^2\n")

for i in range(8):
    target = np.c_[np.ones((X_test.shape[0], 1)), X_test[:,i]]
    y_pred = predict(target, thetas[i])
    MSE = mse(y_test, y_pred)
    var = np.var(y_test)
    rsquared = 1 - (MSE/var)
    print(f"{rsquared}, {MSE}, {var} = R^2 for Unnormalized Feature_
↪{shortenedCols[i]}")

print("")

for i in range(8):
    target = np.c_[np.ones((X_test_norm.shape[0], 1)), X_test_norm[:,i]]
    y_pred = predict(target, norm_thetas[i])
    MSE = mse(y_test_norm, y_pred)
    var = np.var(y_test_norm)
    rsquared = 1 - (MSE/var)
    print(f"{rsquared}, {MSE}, {var} = R^2 for Normalized Feature_
↪{shortenedCols[i]}")

```

Performance on Testing Dataset

RSquared	MSE	Variance	
0.5537961158499328,	114.75321894547034,	257.1766473168498	= Unnormalized Multivariate R ²
0.5717923301901839,	0.39151032813045994,	0.9143001298980588	= Normalized Multivariate R ²
0.05295064793759763,	243.5589772070036,	257.1766473168498	= R ² for Unnormalized Feature Cement
-2.9789647029597193,	1023.2968020992658,	257.1766473168498	= R ² for Unnormalized Feature Slag
-3.0935869602781816,	1052.7749699443173,	257.1766473168498	= R ² for Unnormalized Feature Ash
-0.26953186303719656,	326.4939481978204,	257.1766473168498	= R ² for Unnormalized Feature Water
-4.584213974528381,	1436.1294278691096,	257.1766473168498	= R ² for Unnormalized Feature Superplasticizer
-0.08806069419969553,	279.82380141152186,	257.1766473168498	= R ² for Unnormalized Feature Coarse
-0.1341197658079767,	291.66911902626634,	257.1766473168498	= R ² for Unnormalized Feature Fine

-2.6874385052329544, 948.3230719628674, 257.1766473168498 = R^2 for Unnormalized Feature Age

0.19548643481707606, 0.7355668571514977, 0.9143001298980588 = R^2 for Normalized Feature Cement

-0.027999295268484214, 0.9398998891990881, 0.9143001298980588 = R^2 for Normalized Feature Slag

-0.021315368261298984, 0.9337887738681894, 0.9143001298980588 = R^2 for Normalized Feature Ash

0.07601524974208251, 0.8447993771846394, 0.9143001298980588 = R^2 for Normalized Feature Water

0.1165907335861569, 0.8077012070353256, 0.9143001298980588 = R^2 for Normalized Feature Superplasticizer

-0.008560540991194188, 0.9221270336383053, 0.9143001298980588 = R^2 for Normalized Feature Coarse

-0.01346193263703599, 0.9266083766567798, 0.9143001298980588 = R^2 for Normalized Feature Fine

0.08400049819338729, 0.8374984634883431, 0.9143001298980588 = R^2 for Normalized Feature Age

```
[ ]: # Plots of your univariate models on top of scatterplots of the training data
# used. Please plot the data using the x-axis for the predictor variable and
# the y-axis
# for the response variable.

for i in range(8):
    target = np.array(X_train[:, i])
    plt.figure(figsize=(8, 6))
    plt.scatter(target, y_train, s=10)
    plt.plot(target, np.dot(np.c_[np.ones(target.shape), target], thetas[i].T),
             color='r')
    plt.xlabel(data.columns[i])
    plt.ylabel('Concrete Strength')
    plt.title(data.columns[i] + " vs Concrete Strength")
    plt.show()
```
















